

Neural Network Step by Step - Part I

by Timothy (tq.wang93@hotmail.com)


First draft: Dec, 2020

Last update: Dec, 2020


Logistic regression from scratch

In part I, we will work on the simplest neural network, which includes an **input layer** with n nodes (each representing a **feature** of the input) and **one** node in the **output layer**. This type of neural network can be used to identify whether the item belongs to a certain category (Yes = 1; No = 0).

For example, given a series of features (x_1, x_2, \dots, x_n) of a picture, the model can judge whether it is a cat (Yes = 1; No = 0).



We will train this neural network using m samples. Before showing codes, let's go through the calculation of **feed-forward** and **backward propagation** process.



Feed-forward

- **For sample i**

Let's first define the vector representation of sample i and weight.

$$x^{(i)} = \begin{pmatrix} x_1^i \\ x_2^i \\ \vdots \\ x_n^i \end{pmatrix}_{n \times 1} \quad w = \begin{pmatrix} w_1 \\ w_2 \\ \vdots \\ w_n \end{pmatrix}_{n \times 1}$$

We can then calculate z for sample i and represent it in a more concise way.

$$\begin{aligned} z^{(i)} &= x_1^i w_1 + x_2^i w_2 + \dots + x_n^i w_n + b_i \\ &= (w_1, w_2, \dots, w_n) \begin{pmatrix} x_1^i \\ x_2^i \\ \vdots \\ x_n^i \end{pmatrix} + b_i \\ &= w^T x^{(i)} + b_i \end{aligned}$$

- **For all samples**


$$\begin{aligned} Z &= [z^{(1)}, z^{(2)}, \dots, z^{(m)}]_{1 \times m} \\ &= [w^T x^{(1)} + b_1, w^T x^{(2)} + b_2, \dots, w^T x^{(m)} + b_m] \\ &= w^T [x^{(1)}, x^{(2)}, \dots, x^{(m)}] + [b_1, b_2, \dots, b_m] \\ &= w^T X + b \end{aligned}$$

$$w^T = [w_1, w_2, \dots, w_n]_{1 \times n} \quad X = \begin{pmatrix} x_1^1 & x_1^2 & \dots & x_1^m \\ x_2^1 & x_2^2 & \dots & x_2^m \\ \vdots & \vdots & \ddots & \vdots \\ x_n^1 & x_n^2 & \dots & x_n^m \end{pmatrix}_{n \times m}$$

$$\hat{Y} = \text{sigmoid}(Z) = [\hat{y}^{(1)}, \hat{y}^{(2)}, \dots, \hat{y}^{(m)}]_{1 \times m}$$

Back propagation

Back propagation starts by comparing the **predicted value** \hat{y} to the **true value** y , which produces the **loss**. It then updates the weights w and bias b using **gradient descent**. Through iterations, the weights and bias are optimized.



- **Loss function** $L(a, y)$

Here, we use the **cross entropy** loss function. There are other options such as **mean squared error**.

$$L(a, y) = -y \log a - (1 - y) \log(1 - a)$$

- **Cost function** $J(a, y)$

As defined, cost function is the **mean** of the loss of all samples.

$$J(a, y) = J(w, b) = \frac{1}{m} \sum_{i=1}^m L(a, y)$$

- **Gradient descent**

Based on chain rules,

$$\frac{\partial L(a, y)}{\partial w} = \frac{\partial L(a, y)}{\partial a} \cdot \frac{\partial a}{\partial z} \cdot \frac{\partial z}{\partial w}$$

$$\frac{\partial L(a, y)}{\partial b} = \frac{\partial L(a, y)}{\partial a} \cdot \frac{\partial a}{\partial z} \cdot \frac{\partial z}{\partial b}$$

Note: the *sigmoid* function is

$$a = \sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\begin{aligned} \frac{\partial L(a, y)}{\partial a} &= -y \frac{1}{a} - (1 - y) \frac{1}{1 - a} (-1) \\ &= -\frac{y}{a} + \frac{1 - y}{1 - a} \end{aligned}$$

$$\begin{aligned} \frac{\partial a}{\partial z} &= \frac{0 - 1 \cdot e^{-z} \cdot (-1)}{(1 + e^{-z})^2} \\ &= \frac{e^{-z}}{(1 + e^{-z})^2} \\ &= \frac{1}{1 + e^{-z}} \cdot \frac{e^{-z}}{1 + e^{-z}} \\ &= \frac{1}{1 + e^{-z}} \cdot \left(1 - \frac{1}{1 + e^{-z}}\right) \\ &= a \cdot (1 - a) \end{aligned}$$

Therefore,

$$\begin{aligned} \frac{\partial L(a, y)}{\partial a} \cdot \frac{\partial a}{\partial z} &= \left(-\frac{y}{a} + \frac{1 - y}{1 - a}\right) a(1 - a) \\ &= -y(1 - a) + a(1 - y) \\ &= a - y \end{aligned}$$

Moving on, for the partial derivatives of $z = w^T + b$ to w and b ,

- In the case of **sample i**, because

$$\frac{\partial z}{\partial w} = x_{n \times 1} \quad \frac{\partial z}{\partial b} = 1_{1 \times 1}$$

we can get

$$\frac{\partial L(a, y)}{\partial w} = x_{n \times 1} (a - y)_{1 \times 1}$$

$$\frac{\partial L(a, y)}{\partial b} = (a - y)_{1 \times 1} \cdot 1_{1 \times 1}$$

- In the case of **m samples**,

$$\frac{\partial L(A, Y)}{\partial w_{n \times 1}} = X_{n \times m} (A - Y)^T_{m \times 1}$$

$$\begin{aligned} \frac{\partial L(A, Y)}{\partial b_{1 \times 1}} &= A - Y_{1 \times m} \\ &= np.sum(A - Y)_{1 \times 1} \end{aligned}$$

- **Updating weights and bias**

The weight w and bias b are updated by subtracting the **mean of cost** multiplied by the **learning rate** α .

$$w \leftarrow w - \alpha \cdot \frac{1}{m} \cdot X(A - Y)^T$$

$$b \leftarrow b - \alpha \cdot \frac{1}{m} \cdot np.sum(A - Y)$$

- **Revisiting the cost function J**

Usually, we print/plot the change of cost as a function of iterations.

$$\begin{aligned} J &= \frac{1}{m} \sum_{i=1}^m [-y \log a - (1 - y) \log(1 - a)] \\ &= -\frac{1}{m} \sum_{i=1}^m [y \log a + (1 - y) \log(1 - a)] \\ &= -\frac{1}{m} [Y \log A + (1 - Y) \log(1 - A)] \\ &= -\frac{1}{m} np.sum[Y \log A + (1 - Y) \log(1 - A)] \end{aligned}$$

Note: the $Y \log A + (1 - Y) \log(1 - A)$ is calculated using dot product.

Coding a logistic regression

Code: sigmoid function

$$a = \frac{1}{1 + e^{-z}}$$

```
# define sigmoid function
def sigmoid(z):
    a = 1 / (1 + np.exp(-z))
    return a
```

Code: feed-forward, cost function and gradient descent

- **Feed-forward**

$$Z = w^T X + b$$

$$A = \sigma(Z)$$

- **Cost function**

$$J = -\frac{1}{m} np.sum[Y \log A + (1 - Y) \log(1 - A)]$$

- **Gradient descent**

$$\frac{\partial L(A, Y)}{\partial w} = X(A - Y)^T$$

$$\frac{\partial L(A, Y)}{\partial b} = np.sum(A - Y)$$

```
# initialize parameters
n_dim = train_data_sta.shape[0] # number of rows in training data
w = np.zeros((n_dim, 1))
b = 0

# propagate
def propagate(w, b, x, y):

    # feed-forward function
    Z = np.dot(w.T, x) + b # np.dot -> matrix multiplication
    A = sigmoid(Z)

    # cost function
    m = x.shape[1]
    J = -1/m * np.sum(y * np.log(A) + (1-y) * np.log(1-A))

    # gradient descent (Note: mean)
    dw = 1/m * np.dot(x, (A-y).T)
    db = 1/m * np.sum(A-y)

    grands = {'dw': dw, 'db': db}

    return grands, J
```

Code: optimization

$$w \leftarrow w - \alpha \cdot \frac{1}{m} \cdot X(A - Y)^T$$

$$b \leftarrow b - \alpha \cdot \frac{1}{m} \cdot np.sum(A - Y)$$

```
# Optimization
def optimize(w, b, x, y, alpha, n_iters):
    costs = []
    for i in range(n_iters):
        grands, J = propagate(w, b, x, y)
        dw = grands['dw']
        db = grands['db']

        w = w - alpha * dw
        b = b - alpha * db

        if i % 100 == 0:
            costs.append(J)
            print('Epoch %d: cost = %.4f' % (i+1, J))

    grands = {'dw': dw, 'db': db}
    params = {'w': w, 'b': b}

    return grands, params, costs
```

Code: prediction

```
# Prediction
def predict(w, b, X_test):

    Z = np.dot(w.T, X_test) + b
    A = sigmoid(Z)

    m = X_test.shape[1]
    Y_pred = np.zeros((1, m))

    for i in range(m):
        if A[:, i] > 0.5:
            Y_pred[:, i] = 1
        else:
            Y_pred[:, i] = 0

    return Y_pred
```

Code: integrating previous steps

```
# integrating previous steps
def model(w, b, X_train, X_test, Y_train, Y_test, alpha, n_iters):
    grands, params, costs = optimize(w, b, X_train, Y_train, alpha, n_iters)
    w = params['w']
    b = params['b']

    Y_pred_train = predict(w, b, X_train)
    Y_pred_test = predict(w, b, X_test)

    print('Train accuracy: %.2f' % np.mean(y_pred_train == y_train))
    print('Test accuracy: %.2f' % np.mean(y_pred_test == y_test))

    dic = {
        'w': w,
        'b': b,
        'costs': costs,
        'y_pred_train': y_pred_train,
        'y_pred_test': y_pred_test,
        'alpha': alpha
    }

    return dic
```

Code: train and test

- Train and test

```
dic = model(w, b,
            train_data_sta, train_labels_tran,
            test_data_sta, test_labels_tran,
            alpha = 0.005, n_iters = 2000
            )
```

- Plot the change of cost as a function of iterations

```
plt.plot(b['costs'])  
plt.xlabel('per hundred iterations')  
plt.ylabel('cost')
```

Code: predict a picture

```
index = 1  
print('True label: %d' % test_labels_tran[0, index])  
print('Pred label: %d' % int(b['y_pred_test'][0, index]))  
  
# show the picture  
plt.imshow(test_data_org[index])
```